

Titanic Survivor Classification

Description

This is notebook to build classification for Titanic passengers to classify if those passengers survived or not. Our aim is to get model with the most accuracy possible.

source : <https://www.kaggle.com/competitions/titanic/data>

Data Overview

```
In [1]: import pandas as pd
data = pd.read_csv('train.csv')
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass         891 non-null   int64
3   Name           891 non-null   object
4   Sex            891 non-null   object
5   Age            714 non-null   float64
6   SibSp          891 non-null   int64
7   Parch          891 non-null   int64
8   Ticket         891 non-null   object
9   Fare           891 non-null   float64
10  Cabin          204 non-null   object
11  Embarked       889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
In [2]: data.describe()
```

```
Out[2]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```
In [3]: data.head()
```

Out[3]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

In [4]: `data.nunique()`

Out[4]:

PassengerId	891
Survived	2
Pclass	3
Name	891
Sex	2
Age	88
SibSp	7
Parch	7
Ticket	681
Fare	248
Cabin	147
Embarked	3
dtype:	int64

So, here is Describe for each column

- PassengerId : id for each passenger for this data
- Survived : if this passenger survived or not (0 = No, 1 = Yes)
- Pclass : Ticket class (1 = 1st, 2 = 2nd, 3 = 3rd)
- Name : name of passenger
- Sex : sex of passenger
- Age : age of passenger
- SibSp : number of siblings / spouses aboard the Titanic
- Parch : number of parents / children aboard the Titanic
- Ticket : Ticket number
- Fare : Passenger fare
- Cabin : Cabin number
- Embarked : Port of Embarkation (C = Cherbourg, Q = Queenstown, S = Southampton)

Clean Up

Remove column that we are not going to use and fill null data.

```
In [5]: # remove cloumns PassengerId, Name, Ticket and Cabin Since it have high cardinality.
data = data[['Survived', 'Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Embarked']]
```

```
In [6]: # get mean age
mean_age = data['Age'].mean()

print('mean age = ', mean_age)

# show histogram for Embarked for impute with the most frequent value
mode_emb = data['Embarked'].mode()
print('mode Embarked = ', mode_emb[0])
```

```
mean age = 29.69911764705882
mode Embarked = S
```

```
In [7]: # impute missing data
values = {'Age' : mean_age, 'Embarked': 'S'}
data = data.fillna(value=values)
```

```
In [8]: # convert Sex to int and Embarked to 3 int columns
data['Sex'] = data.Sex.apply(lambda x: int(x == 'male'))
data['Emb_C'] = data.Embarked.apply(lambda x: int(x == 'C'))
data['Emb_Q'] = data.Embarked.apply(lambda x: int(x == 'Q'))
data['Emb_S'] = data.Embarked.apply(lambda x: int(x == 'S'))

data = data[['Survived', 'Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Emb_C', 'Emb_Q', 'Emb_S']]
data.head()
```

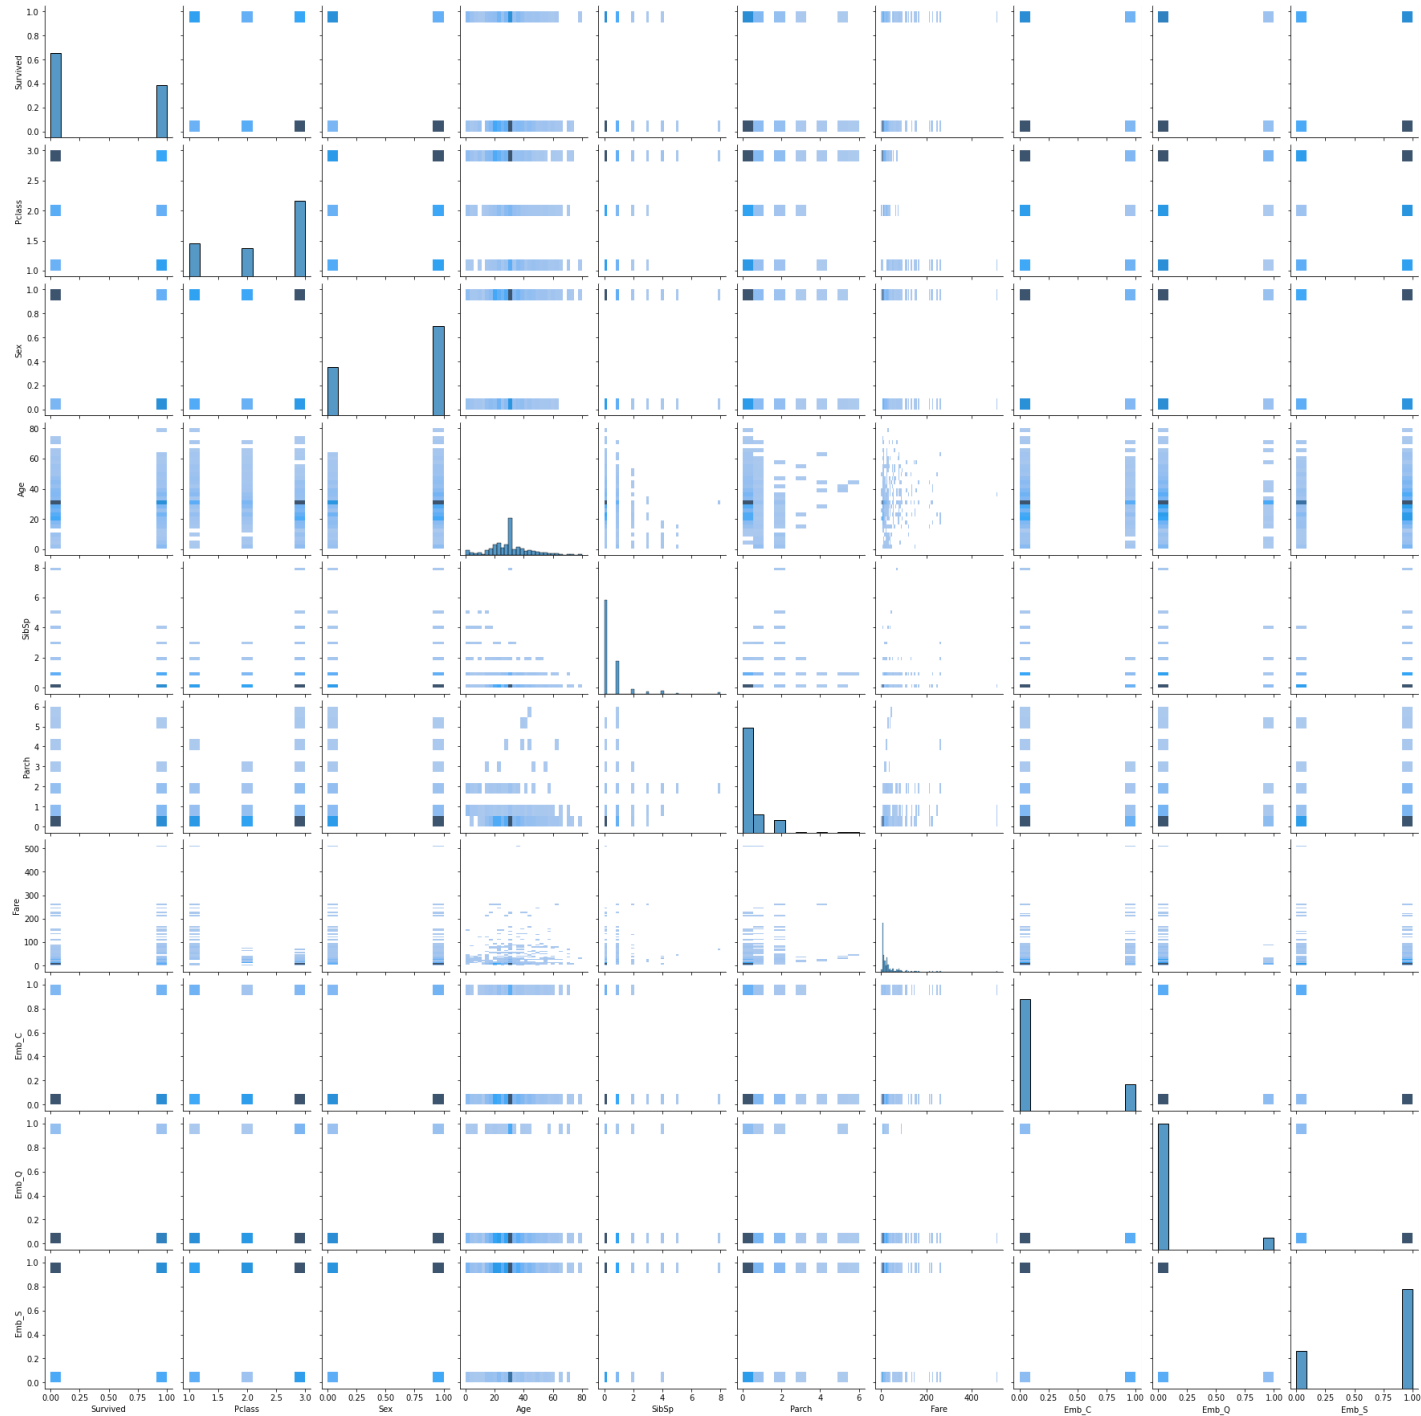
```
Out[8]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Emb_C	Emb_Q	Emb_S
0	0	3	1	22.0	1	0	7.2500	0	0	1
1	1	1	0	38.0	1	0	71.2833	1	0	0
2	1	3	0	26.0	0	0	7.9250	0	0	1
3	1	1	0	35.0	1	0	53.1000	0	0	1
4	0	3	1	35.0	0	0	8.0500	0	0	1

Visual Display

```
In [9]: import seaborn as sns
import matplotlib.pyplot as plt

sns.pairplot(data, kind="hist")
plt.savefig('pair_plot1.png')
```



In [10]:

```
corr_data = data.corr()
print(corr_data)
```

	Survived	Pclass	Sex	Age	SibSp	Parch	\
Survived	1.000000	-0.338481	-0.543351	-0.069809	-0.035322	0.081629	
Pclass	-0.338481	1.000000	0.131900	-0.331339	0.083081	0.018443	
Sex	-0.543351	0.131900	1.000000	0.084153	-0.114631	-0.245489	
Age	-0.069809	-0.331339	0.084153	1.000000	-0.232625	-0.179191	
SibSp	-0.035322	0.083081	-0.114631	-0.232625	1.000000	0.414838	
Parch	0.081629	0.018443	-0.245489	-0.179191	0.414838	1.000000	
Fare	0.257307	-0.549500	-0.182333	0.091566	0.159651	0.216225	
Emb_C	0.168240	-0.243292	-0.082853	0.032024	-0.059528	-0.011069	
Emb_Q	0.003650	0.221009	-0.074115	-0.013855	-0.026354	-0.081228	
Emb_S	-0.149683	0.074053	0.119224	-0.019336	0.068734	0.060814	

	Fare	Emb_C	Emb_Q	Emb_S
Survived	0.257307	0.168240	0.003650	-0.149683
Pclass	-0.549500	-0.243292	0.221009	0.074053
Sex	-0.182333	-0.082853	-0.074115	0.119224


```

Age      0.091566  0.032024 -0.013855 -0.019336
SibSp    0.159651 -0.059528 -0.026354  0.068734
Parch    0.216225 -0.011069 -0.081228  0.060814
Fare     1.000000  0.269335 -0.117216 -0.162184
Emb_C    0.269335  1.000000 -0.148258 -0.782742
Emb_Q   -0.117216 -0.148258  1.000000 -0.499421
Emb_S   -0.162184 -0.782742 -0.499421  1.000000

```

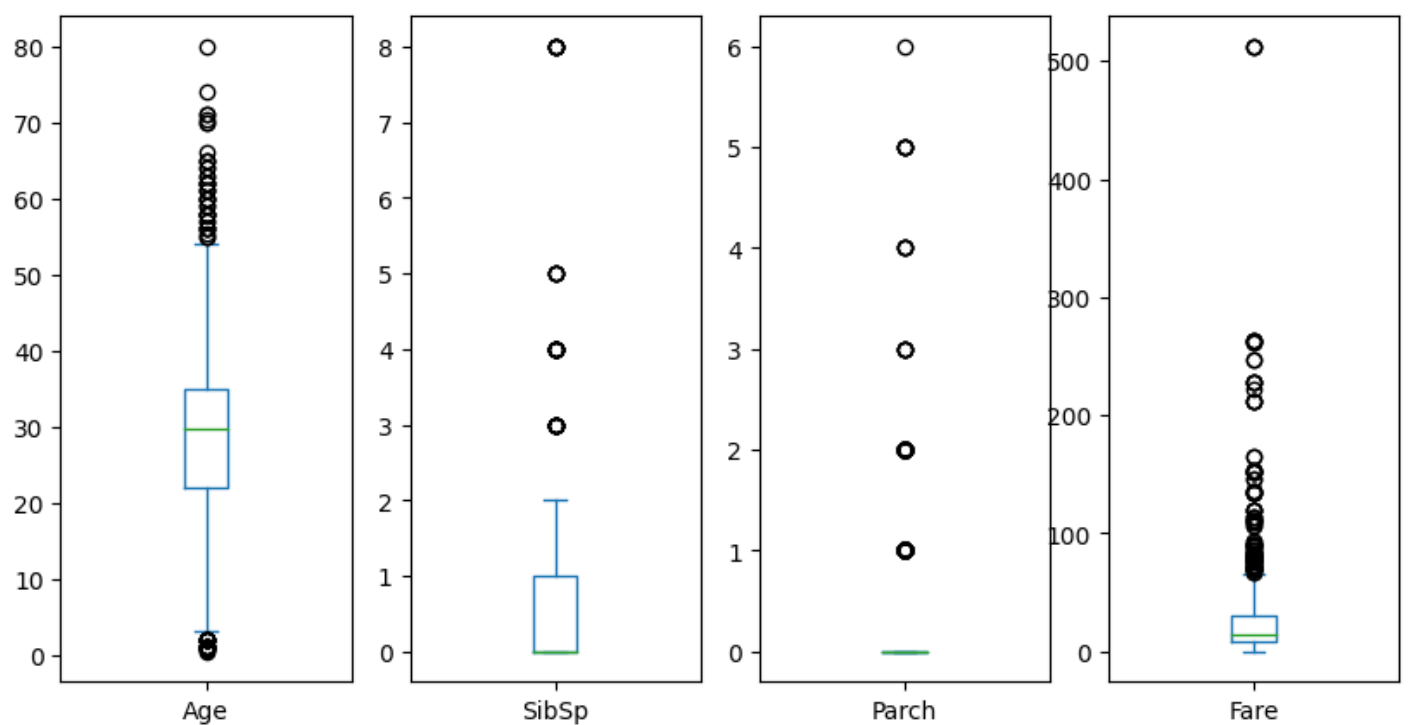
```
In [11]: sns.heatmap(corr_data)
```

```
Out[11]: <AxesSubplot:>
```



```
In [12]: plot_data = data[['Age', 'SibSp', 'Parch', 'Fare']]
plot_data.plot(kind="box", subplots=True, figsize=(10, 5))
```

```
Out[12]: Age      AxesSubplot(0.125,0.11;0.168478x0.77)
SibSp    AxesSubplot(0.327174,0.11;0.168478x0.77)
Parch    AxesSubplot(0.529348,0.11;0.168478x0.77)
Fare     AxesSubplot(0.731522,0.11;0.168478x0.77)
dtype: object
```



Analysis

From information above, seems there is no strong correlation between features as you can see from correlation heatmap or raw data (>0.7). You may see Emb_C and Emb_S have high value but those columns were extracted from Embarked feature. There are a lot of outlier data as you can see. But, I think those are all valid. Age is in a reasonable range (more than 0 and less than 90). SibSp and Parch are also reasonable because most of people board alone. Fare is also reasonable for me.

Classification

I want to use Random Forest Tree to solve this problem. Because we have only a few features.

First, split data into train and test (test for 20%).

```
In [23]: from sklearn.model_selection import train_test_split

X = data[['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Emb_C', 'Emb_Q', 'Emb_S']]
Y = data['Survived']
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=1992)

print(len(x_train), len(x_test))
```

712 179

Then, Build first model to be our base line.

```
In [24]: from sklearn.ensemble import RandomForestClassifier

bl_model = RandomForestClassifier(n_estimators=10, max_depth=5, random_state=1992).fit(x_train, y_train)
print('score with training data = ', bl_model.score(x_train, y_train))
print('score with test data = ', bl_model.score(x_test, y_test))
```

score with training data = 0.8539325842696629
score with test data = 0.7653631284916201

try to play with n_estimators and max_depth

```

In [25]: es = list(range(10,110,10))
train_scores = []
test_scores = []
for e in es:
    model = RandomForestClassifier(n_estimators=e, max_depth=5, random_state=1992).fit(x_train, y_train)
    train_scores.append(model.score(x_train, y_train))
    test_scores.append(model.score(x_test, y_test))

print('train_scores = ', train_scores)
print('test_scores = ', test_scores)

plt.plot(es, train_scores)
plt.plot(es, test_scores)

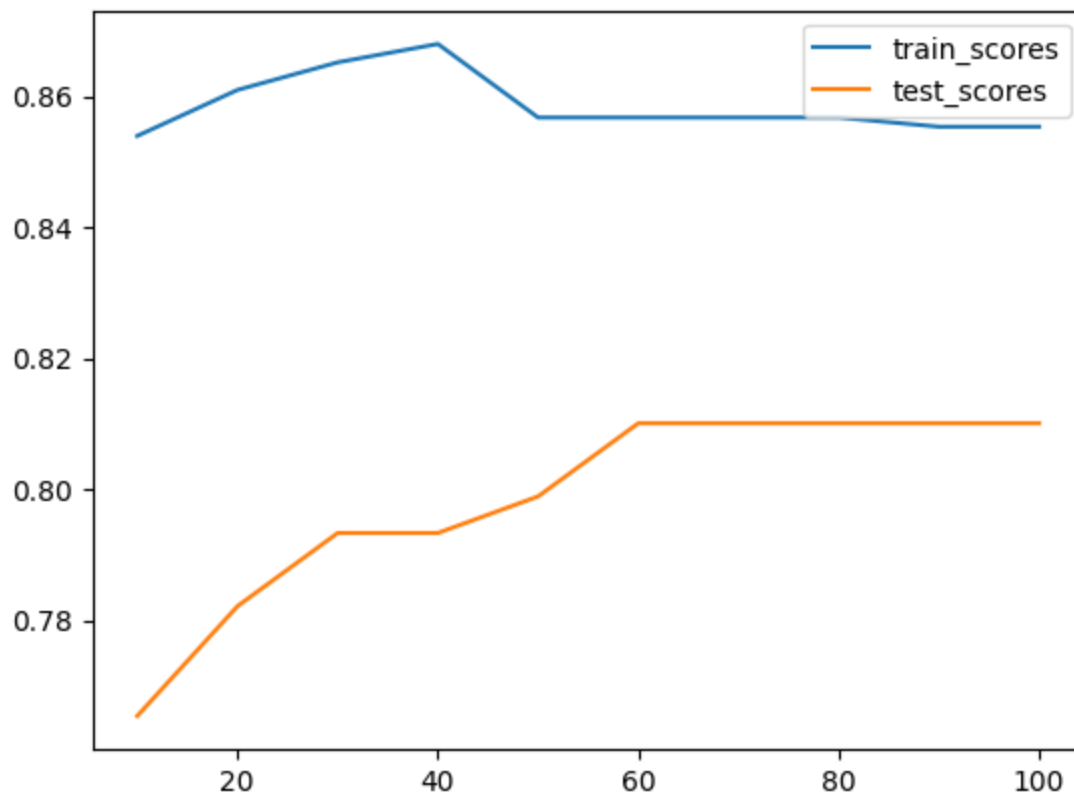
plt.legend(['train_scores', 'test_scores'])
plt.show()

```

```

train_scores = [0.8539325842696629, 0.8609550561797753, 0.8651685393258427, 0.86797752808
98876, 0.8567415730337079, 0.8567415730337079, 0.8567415730337079, 0.8567415730337079, 0.8
553370786516854, 0.8553370786516854]
test_scores = [0.7653631284916201, 0.7821229050279329, 0.7932960893854749, 0.793296089385
4749, 0.7988826815642458, 0.8100558659217877, 0.8100558659217877, 0.8100558659217877, 0.81
00558659217877, 0.8100558659217877]

```



Try with max_depth=10

```

In [28]: es = list(range(10,110,10))
train_scores = []
test_scores = []
for e in es:
    model = RandomForestClassifier(n_estimators=e, max_depth=10, random_state=1992).fit(x_train, y_train)
    train_scores.append(model.score(x_train, y_train))
    test_scores.append(model.score(x_test, y_test))

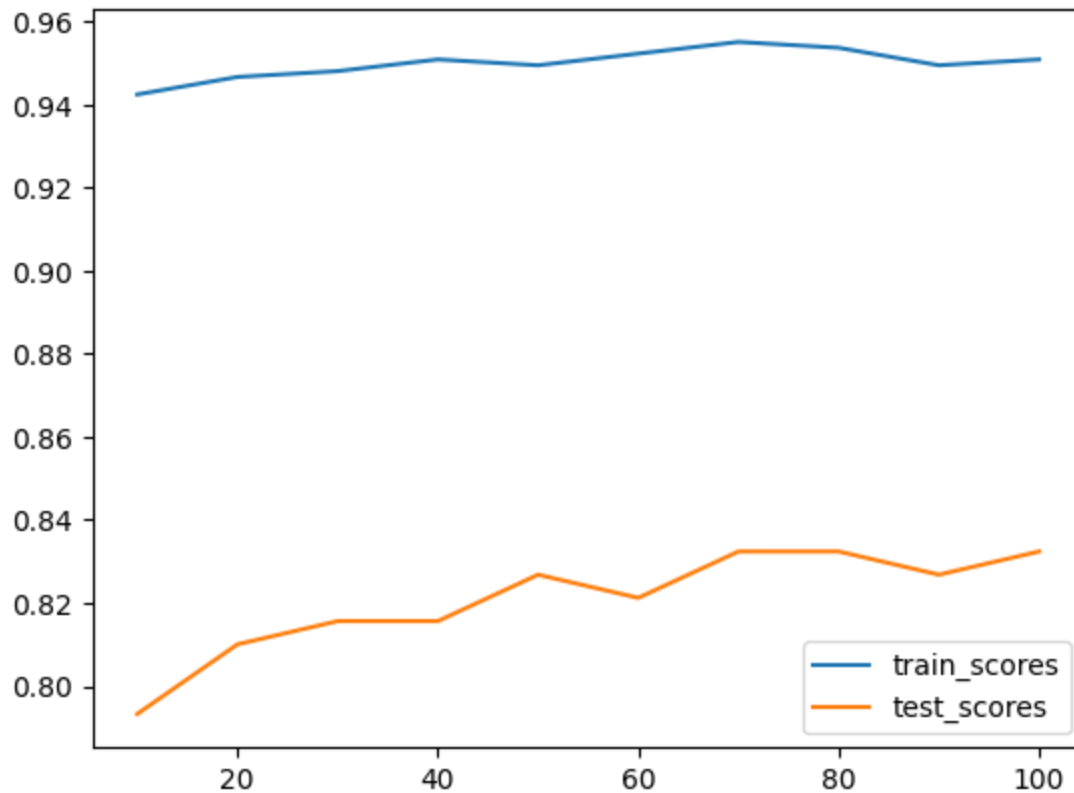
print('train_scores = ', train_scores)
print('test_scores = ', test_scores)

plt.plot(es, train_scores)
plt.plot(es, test_scores)

```

```
plt.legend(['train_scores', 'test_scores'])
plt.show()
```

```
train_scores = [0.9424157303370787, 0.9466292134831461, 0.9480337078651685, 0.95084269662
92135, 0.949438202247191, 0.952247191011236, 0.9550561797752809, 0.9536516853932584, 0.949
438202247191, 0.9508426966292135]
test_scores = [0.7932960893854749, 0.8100558659217877, 0.8156424581005587, 0.815642458100
5587, 0.8268156424581006, 0.8212290502793296, 0.8324022346368715, 0.8324022346368715, 0.82
68156424581006, 0.8324022346368715]
```



From graphs we can see that using max_depth=10 is better. And the reasonable n_estimators is 50.

I will try to tune ccp_alpha to reduce data overfitted.

In [29]:

```
alphas = []
for i in range(10):
    alphas.append(i*0.005)

train_scores = []
test_scores = []
for a in alphas:
    model = RandomForestClassifier(n_estimators=50, max_depth=10, ccp_alpha=a, random_state=42)
    train_scores.append(model.score(x_train, y_train))
    test_scores.append(model.score(x_test, y_test))


print('train_scores = ', train_scores)
print('test_scores = ', test_scores)

plt.plot(alphas, train_scores)
plt.plot(alphas, test_scores)

plt.legend(['train_scores', 'test_scores'])
plt.show()
```

```
train_scores = [0.949438202247191, 0.8595505617977528, 0.8356741573033708, 0.814606741573
0337, 0.8188202247191011, 0.7921348314606742, 0.7879213483146067, 0.7879213483146067, 0.78
79213483146067, 0.7879213483146067]
test_scores = [0.8268156424581006, 0.7988826815642458, 0.7988826815642458, 0.787709497206
```


7039, 0.770949720670391, 0.7821229050279329, 0.7821229050279329, 0.7821229050279329, 0.7821229050279329, 0.7821229050279329]



Increase `ccp_alpha` helps reduce over fitted problem. But, I don't want to sacrifice score for this.

So, the best model for me is `n_estimators=50`, `max_depth=10` and `ccp_alpha=0`

```
In [36]: best_model = RandomForestClassifier(n_estimators=50, max_depth=10, ccp_alpha=0, random_state=42)
print('score with training data = ',bl_model.score(x_train,y_train))
print('score with test data = ',bl_model.score(x_test,y_test))
```

score with training data = 0.8539325842696629

score with test data = 0.7653631284916201

Use this setting and train with all data we have to predict test data from Kaggle.

```
In [31]: best_model = RandomForestClassifier(n_estimators=50, max_depth=10, ccp_alpha=0, random_state=42)
print('score with training data = ',bl_model.score(X,Y))
```

score with training data = 0.8361391694725028

```
In [32]: org_test_data = pd.read_csv('test.csv')
test_data = org_test_data[['Pclass','Sex','Age','SibSp','Parch','Fare','Embarked']]

mean_fare = data['Fare'].mean()

values = {'Age' : mean_age, 'Fare': mean_fare, 'Embarked': 'S'}
test_data = test_data.fillna(value=values)
```

```
In [33]: test_data['Sex'] = test_data.Sex.apply(lambda x: int(x == 'male'))
test_data['Emb_C'] = test_data.Embarked.apply(lambda x: int(x == 'C'))
test_data['Emb_Q'] = test_data.Embarked.apply(lambda x: int(x == 'Q'))
test_data['Emb_S'] = test_data.Embarked.apply(lambda x: int(x == 'S'))

test_data = test_data[['Pclass','Sex','Age','SibSp','Parch','Fare','Emb_C', 'Emb_Q', 'Emb_S']]
```

```
In [34]: predictions = best_model.predict(test_data)
output = pd.DataFrame({'PassengerId': org_test_data.PassengerId, 'Survived': predictions})
output.to_csv('submission.csv', index=False)
print("Your submission was successfully saved!")
```

Your submission was successfully saved!

Score for submit this prediction to Kaggle = 0.77751

Conclusion

For this problem first we look at overview of data we have. Then, we clean it up (select only useful feature fill null data). After that we do analysis about correlation and outlier. And, we build classification model with random forest tree and tune it to get the best model. Last, we use our best model to do prediction and submit it to Kaggle.

Actually, I also try Adaboost but it worst than this. But I still think that there are better approach to this problem which I haven't try it.

github : <https://github.com/Satjarporn/Titanic>

In []: